



VIT[®]
Vellore Institute of Technology
(Deemed to be University under section 3 of UGC Act, 1956)

WINTER SEMESTER 2026-2027

Project Title:
End-to-End Visual Speech Recognition
On the GRID Corpus

Deep Learning (PAMCA605)

Submitted By

Abhay Verma

Reg No: 25MCA1070

Under the guidance of:

Dr. Saleena B

**Assistant Professor in School of Computer Science Engineering
(SCOPE)**

Vellore Institute of Technology, Chennai.

TABLE OF CONTENTS

Abstract	iv
1. Introduction.....	1
1.1 Problem Background	1
1.2 Objectives	2
2. Dataset Description.....	3
2.1 Source and Overview.....	3
2.2 Data Structure and Organization.....	3
2.3 Sentence Grammar and Vocabulary	4
3. Data Preprocessing	5
3.1 Frame Extraction and Validation.....	5
3.2 Mouth Region of Interest Extraction.....	6
3.3 Grayscale Conversion and Normalization.....	6
3.4 Alignment Parsing and Tokenization	6
3.5 TensorFlow Data Pipeline Construction	7
4. Methodology.....	8
4.1 System Overview.....	8
4.2 Algorithm 1: 3D Convolutional Neural Network + Bidirectional LSTM.....	8
4.3 Algorithm 2: 2D Convolutional Neural Network + Transformer Encoder.....	10
4.4 Connectionist Temporal Classification (CTC).....	12
5. Implementation Details	13
5.1 Development Environment and Libraries	13
5.2 Vocabulary Construction and Encoding	14
5.3 Model Hyperparameters and Training Configuration.....	14
5.4 Transfer Learning and Weight Initialization	15
6. Results & Evaluation	16
5.1 Quantitative Performance Metrics	16
5.2 Training and Validation Loss Curves	17
5.3 Qualitative Analysis of Sample Predictions.....	18
7. Comparative Analysis.....	19
7.1 Architectural Trade-offs	19
7.2 Convergence Behavior and Data Efficiency.....	20
7.3 Computational Characteristics	20
8. Conclusion	21

9. Future Work	22
10. References	23

LIST OF FIGURES

Figure	Description	Page
1	High-Level System Architecture of the End-to-End Visual Speech Recognition Pipeline	8
2	Mouth Region of Interest Extraction: Original Frame with Bounding Box vs. Cropped Grayscale Input	6
3	Architecture Diagram of the 3D CNN + Bidirectional LSTM Model	9
4	Architecture Diagram of the 2D CNN + Transformer Encoder Model	11
5	Illustration of Connectionist Temporal Classification (CTC) Alignment and Greedy Decoding	12
6	Training and Validation CTC Loss Curves: BiLSTM Model vs. Transformer Model	17
7	Model Performance Comparison on Test Set: Character Error Rate (CER) and Word Error Rate (WER) Bar Chart	16
8	Sample Prediction Visualization: Ground Truth vs. Model Predictions on Six Test Videos	18

LIST OF TABLES

Table	Description	Page
1	GRID Corpus Sentence Grammar Structure	4
2	Detailed Preprocessing Transformations and Output Shapes	7
3	Model Hyperparameter Summary	14
4	Final Performance Metrics on Held-Out Test Set	16
5	Comparative Analysis of Architectural Characteristics	19

Abstract

This project presents a deep learning-based end-to-end visual speech recognition (VSR) system designed to transcribe spoken language directly from video sequences of facial movements, without any audio signal. VSR—commonly known as lip reading—addresses critical challenges in accessibility for individuals with hearing impairments and enables robust human-computer interaction in acoustically constrained environments.

The primary objective is to develop and rigorously compare two advanced deep learning architectures beyond the prescribed syllabus, utilizing a real-world dataset of approximately 1,000 video samples from the GRID Corpus (Speaker 1 subset). The entire implementation is executed on the Kaggle cloud GPU platform using TensorFlow 2.19.0 with dual NVIDIA T4 GPU acceleration.

The first architecture is a **3D CNN + Bidirectional LSTM (BiLSTM)** hybrid, wherein three-dimensional convolutional layers extract spatiotemporal features from raw video frames and BiLSTM layers model long-range sequential dependencies. This model is fine-tuned from pre-trained checkpoint weights (epoch 96) using Connectionist Temporal Classification (CTC) loss, achieving a **Character Error Rate (CER) of 0.6224** and a **Word Error Rate (WER) of 0.8833** on the held-out test set.

The second architecture is a **CNN + Transformer Encoder** hybrid, wherein a two-dimensional CNN extracts per-frame spatial features and a Transformer Encoder with multi-head self-attention models global temporal relationships across all 75 frames in parallel. Trained from scratch for 20 epochs, this model achieves a **CER of 0.8125** and a **WER of 1.0000**, reflecting the greater data and training epoch requirements characteristic of attention-based architectures.

Comprehensive preprocessing includes mouth Region of Interest (ROI) extraction, grayscale conversion, per-video normalization, and an optimized tf.data pipeline. The comparative analysis demonstrates that while the BiLSTM model benefits from transfer learning and converges rapidly, the Transformer architecture is inherently more scalable and parallelizable, positioning it as a promising long-term direction for visual speech recognition research.

1. INTRODUCTION

1.1 Problem Background

Visual Speech Recognition (VSR) is the task of understanding spoken language from visual information alone—specifically, the dynamic movements of a speaker's lips, teeth, and tongue as captured in video sequences. Unlike conventional automatic speech recognition (ASR) systems that depend on acoustic signals, VSR operates entirely on visual data. This fundamental distinction makes VSR uniquely valuable in several critical application domains:

- **Assistive Technology:** For the millions of individuals worldwide with hearing impairments, robust lip reading technology can provide real-time transcription of spoken conversations, enhancing accessibility and social inclusion.
- **Acoustically Challenged Environments:** In noisy industrial settings, crowded public spaces, or secure facilities where audio recording is prohibited, VSR enables silent speech interfaces.
- **Multimodal Authentication:** Visual speech patterns can serve as an additional biometric factor for identity verification.

Despite its promise, VSR presents formidable computational challenges. Viseme ambiguity is inherent—multiple phonemes produce visually identical lip configurations. For instance, the bilabial consonants /p/, /b/, and /m/ are indistinguishable from visual information alone. Speaker variability introduces significant variance in lip morphology, speaking rate, and articulatory style. Furthermore, the sequence-to-sequence alignment problem is non-trivial: a three-second video comprising 75 frames must be mapped to a variable-length text transcript without explicit segmentation labels.

The field of VSR has been transformed by deep learning. Early approaches relied on hand-crafted visual features coupled with Hidden Markov Models (HMMs), which suffered from limited scalability and required complex, multi-stage training pipelines. The introduction of end-to-end trainable neural networks demonstrated that models could learn directly from raw video frames, bypassing intermediate phoneme classification. This project builds upon these advances, implementing and critically comparing two architecturally distinct deep learning paradigms for sentence-level visual speech recognition on the GRID Corpus.

1.2 Objectives

The development of this project was guided by the following primary objectives:

- **Data Acquisition and Preprocessing:** Download the GRID Corpus (Speaker 1 subset), extract mouth Region of Interest (ROI) from video frames, convert to grayscale, apply per-video normalization, and construct an efficient tf.data pipeline for GPU-accelerated training.
- **Implementation of Spatiotemporal Convolutional Recurrent Architecture:** Design and implement a 3D CNN + Bidirectional LSTM model. The three-dimensional convolutions capture joint spatial and temporal features, while the BiLSTM layers model long-range sequential dependencies in both forward and backward directions.
- **Implementation of Convolutional Transformer Architecture:** Design and implement a 2D CNN + Transformer Encoder model. A shared two-dimensional CNN extracts per-frame spatial

embeddings, and a Transformer Encoder with multi-head self-attention models global temporal context across the entire sequence.

- **Model Training and Transfer Learning:** Fine-tune the BiLSTM model from pre-trained checkpoint weights (epoch 96) and train the Transformer model from scratch with random initialization.
- **Performance Evaluation:** Evaluate both models using industry-standard metrics—Character Error Rate (CER) and Word Error Rate (WER)—on a 100-sample held-out test set.
- **Comparative Analysis:** Conduct a rigorous comparative analysis evaluating accuracy, convergence speed, data efficiency, computational characteristics, and architectural scalability of both approaches.

2. Dataset Description

2.1 Source and Overview

This project utilizes the **GRID Audio-Visual Corpus** (Cooke et al., 2006), a widely adopted benchmark in audio-visual speech recognition research. The corpus comprises high-quality frontal videos of 34 native English speakers uttering short sentences in a controlled studio environment. The consistent lighting, fixed camera angle, and uniform background make the GRID Corpus particularly suitable for developing and benchmarking VSR models without confounding environmental variables.

For this project, the **Speaker 1 (s1)** subset is employed, providing a focused dataset of exactly 1,000 video samples. The videos are stored in MPEG format (.mpg) with a resolution of 720×576 pixels. Each video is approximately three seconds in duration, recorded at a frame rate of 25 frames per second (fps), yielding a consistent 75 frames per video after clipping. Corresponding word-level alignment files (.align) provide precise start and end frame annotations for each spoken word.

Attribute	Description
Corpus	GRID Audio-Visual Corpus
Speaker	Speaker 1 (s1)
Total Videos	1,000 .mpg files
Frame Rate	25 frames per second (fps)
Duration per Video	Approximately 3 seconds (75 frames)
Alignment Files	1,000 corresponding .align text files
Train / Test Split	900 training samples / 100 test samples
Download Method	Google Drive via gdown utility

2.2 Data Structure and Organization

The dataset follows a clean, parallel directory structure that facilitates efficient file pairing during preprocessing:

- **Video Directory** (data/s1/): Contains 1,000 .mpg video files. Each filename is a structured alphanumeric code that encodes the sentence content. For example, the file bbaf2n.mpg corresponds to the sentence "bin blue at f two now."
- **Alignment Directory** (data/alignments/s1/): Contains 1,000 corresponding .align text files. Each alignment file shares the same base filename as its corresponding video (e.g., bbaf2n.align).

Each .align file contains lines in the following format:

<i>start_frame</i>	<i>end_frame</i>	<i>word</i>
0	12	<i>sil</i>
12	18	<i>bin</i>
18	24	<i>blue</i>
...		

The **sil** tokens denote silent periods at the beginning and end of each recording. These are programmatically excluded from the target transcript during preprocessing to ensure the model learns only the spoken content.

2.3 Sentence Grammar and Vocabulary

Every sentence in the GRID Corpus adheres to a strict six-word command grammar, which significantly constrains the syntactic complexity and vocabulary size. This controlled structure makes the dataset particularly well-suited for training deep learning models with limited samples, as it reduces the combinatorial explosion of possible word sequences.

Position	Category	Options
1	Command	bin, lay, place, set
2	Color	blue, green, red, white
3	Preposition	at, by, in, with
4	Letter	A through Z (excluding W)
5	Digit	zero, one, two, three, four, five, six, seven, eight, nine
6	Adverb	again, now, please, soon

Example Sentence: bin blue at f two now

The restricted vocabulary comprises S1 unique words across all six grammatical categories. This constraint reduces the word-level perplexity of the task, allowing models to focus on learning the visual-to-phonetic mapping rather than complex language modeling.

3. Data Preprocessing

The preprocessing pipeline transforms raw MPEG video files and their corresponding alignment text files into normalized tensor pairs suitable for efficient model training. This section details each stage of the pipeline, from frame extraction to the construction of a high-performance tf.data input pipeline.

3.1 Frame Extraction and Validation


The first stage involves decoding each .mpg video file into a sequence of individual frames. The OpenCV library (cv2.VideoCapture) is employed for this purpose. Each video is opened, and frames are read sequentially until the maximum frame count (75) is reached or the video stream ends.

A robust error-handling mechanism is implemented to address minor MPEG stream corruption present in some GRID Corpus files. When decoding errors occur (manifesting as FFMPEG warnings such as "ac-tex damaged" or "Warning MVs not available"), the affected frame is skipped gracefully. Videos that yield fewer than 75 valid frames are zero-padded to maintain a uniform input shape across the entire dataset. This approach ensures that no sample is discarded due to minor data corruption, preserving the full 1,000-sample dataset.

3.2 Mouth Region of Interest Extraction

To focus the model's attention on the most informative visual cues, the mouth region is cropped from each full frame. The bounding box coordinates [190:236, 80:220] were determined empirically based on the consistent camera setup of the GRID Corpus. This crop yields a region of 46 pixels in height and 140 pixels in width, centered precisely on the speaker's mouth.

The use of a fixed bounding box is appropriate for this controlled dataset, where speaker position and head pose remain constant across all videos. In real-world deployments, this static cropping would be replaced with a dynamic facial landmark detector.



Original Frame with ROI Box



Extracted Mouth Crop
(46×140 pixels)

Figure : Visually demonstrates the critical preprocessing step that isolates the speaker's mouth

3.3 Grayscale Conversion and Normalization

Color information is largely redundant for lip reading; luminance contrast provides the primary signal for distinguishing lip contours and intra-oral visibility. Therefore, each cropped mouth frame is converted from the BGR color space (OpenCV default) to a single-channel grayscale image.

Subsequently, per-video normalization is applied. For each video independently, the mean pixel intensity is subtracted, and the result is divided by the standard deviation plus a small epsilon ($1e-6$) to prevent division by zero:

$$\text{normalized_frame} = (\text{frame} - \text{mean}(\text{frame})) / (\text{std}(\text{frame}) + 1e-6)$$

This per-sample standardization centers the pixel value distribution around zero and scales it to unit variance. This normalization strategy mitigates variations in lighting and contrast across different video recordings, improving training stability and convergence speed.

3.4 Alignment Parsing and Tokenization

The corresponding .align file for each video is parsed to extract the ground truth transcript. The parsing logic reads each line, splits it into components, and collects all words except those labeled as sil (silence). The resulting word sequence is joined with spaces to form the target sentence string.

This sentence string is then tokenized at the character level using TensorFlow's StringLookup layer. The vocabulary consists of 40 characters: lowercase letters a through z, digits 1 through 9, space, apostrophe ('), exclamation mark (!), and question mark (?). Each character is mapped to a unique integer index from 0 to 39. The out-of-vocabulary token is reserved for any unseen characters.

The final target for each video is an integer sequence of variable length (up to approximately 40 tokens), which will be padded during batch construction.

3.5 TensorFlow Data Pipeline Construction

An efficient data pipeline is critical for maximizing GPU utilization during training. The raw Python-based preprocessing operations (video decoding, cropping) are computationally intensive and would become a bottleneck if executed synchronously on the main thread. To address this, a tf.data pipeline is constructed with the following stages:

1. **File Listing:** `tf.data.Dataset.list_files` creates a dataset of all .mpg file paths in the `data/sl/` directory.
2. **Shuffling:** `.shuffle(buffer_size=500)` randomizes the order of samples to ensure that each batch contains a diverse mix of sentences, promoting stable gradient updates.
3. **Parallel Mapping:** `.map(lambda x: tf.py_function(load_sample, [x], [tf.float32, tf.int32]), num_parallel_calls=tf.data.AUTOTUNE)` applies the `load_sample` function. The `tf.py_function` wrapper bridges the gap between Python OpenCV operations and the TensorFlow graph execution mode. The `num_parallel_calls` parameter enables multi-threaded preprocessing.


Preprocessing Stage	Operation	Output Shape
1. Frame Extraction	OpenCV VideoCapture decode	(75, H, W, 3)
2. Grayscale + Crop	BGR $\rightarrow$ Grayscale; Crop [190:236, 80:220]	(75, 46, 140, 1)
3. Normalization	Subtract mean, divide by standard deviation	(75, 46, 140, 1)
4. Alignment Parsing	Parse .align, exclude sil, tokenize	(≤ 40 ,) int32
5. tf.data Pipeline	Shuffle $\rightarrow$ Map $\rightarrow$ Padded Batch $\rightarrow$ Prefetch	Batched dataset

4. **Padded Batching:** `.padded_batch(batch_size=2, padded_shapes=([75, 46, 140, 1], [40]))` groups samples into batches of size 2. The video tensors have a fixed shape and do not require padding, but the variable-length integer sequences are padded to a uniform length of 40 tokens.
5. **Prefetching:** `.prefetch(buffer_size=tf.data.AUTOTUNE)` instructs the pipeline to prepare the next batch of data while the current batch is being processed by the GPU. This overlap of data loading and model computation is essential for achieving high GPU utilization.

The batch size is constrained to **2** due to the substantial memory footprint of the 3D video tensors on the 16 GB NVIDIA T4 GPUs.

4. Methodology

This section provides a comprehensive explanation of the two deep learning algorithms selected for comparison, along with the training methodology and loss function employed.




4.1 System Overview

The overall system follows a modular, end-to-end design. A raw MPEG video file enters the preprocessing pipeline and is transformed into a normalized tensor of shape (Batch, 75, 46, 140, 1). This tensor serves as the input to one of two parallel model backbones. Each backbone consists of a **spatial feature extraction** component followed by a **temporal sequence modeling** component. The output logits are passed through a softmax activation to produce a probability distribution over 41 classes (40

characters plus the CTC blank token) at each of the 75 timesteps. Finally, a CTC decoder converts the per-frame probability sequence into a human-readable text transcript.

4.2 Algorithm 1: 3D Convolutional Neural Network + Bidirectional LSTM

The first architecture is a hybrid model that combines three-dimensional convolutions for spatiotemporal feature extraction with bidirectional recurrent layers for sequence modeling. This design is inspired by foundational work in sentence-level lip reading and treats the input video as a true volumetric data structure.



4.2.1 Conceptual Background

Three-Dimensional Convolutions: Unlike standard 2D convolutions that operate independently on each frame, 3D convolutions apply a three-dimensional kernel that slides across both spatial dimensions (height and width) and the temporal dimension (frame index). This allows the model to learn motion-sensitive filters that capture dynamic patterns such as lip rounding, jaw lowering, and the rapid transitions between phonemes. The early layers of the network can thus detect primitive spatiotemporal features that are essential for disambiguating visual speech.

Bidirectional Long Short-Term Memory (BiLSTM): LSTMs are a specialized type of recurrent neural network designed to mitigate the vanishing gradient problem that plagues standard RNNs. Through a system of gating mechanisms (input gate, forget gate, output gate), LSTMs can maintain a memory cell state over long sequences, selectively remembering or forgetting information. A **bidirectional** LSTM runs two independent LSTM layers—one processing the sequence forward in time and the other backward. The outputs are concatenated, providing each timestep with contextual information from both past and future frames. This is critical for lip reading, where the visual manifestation of a phoneme is often influenced by surrounding co-articulatory effects.

4.2.2 Layer-by-Layer Architecture Specification

The model is constructed using the Keras Sequential API. The following table details the exact layer configuration, including output shapes and parameter counts, as implemented in the project notebook.

Layer (Type)	Configuration	Output Shape	Parameters
Input	-	(Batch, 75, 46, 140, 1)	0
Conv3D_1	Filters: 128, Kernel: 3×3×3, Padding: 'same'	(Batch, 75, 46, 140, 128)	3,584
Activation	ReLU	(Batch, 75, 46, 140, 128)	0

Layer (Type)	Configuration	Output Shape	Parameters
MaxPool3D_1	Pool size: (1, 2, 2), Strides: (1, 2, 2)	(Batch, 75, 23, 70, 128)	0
Conv3D_2	Filters: 256, Kernel: 3×3×3, Padding: 'same'	(Batch, 75, 23, 70, 256)	884,992
Activation	ReLU	(Batch, 75, 23, 70, 256)	0
MaxPool3D_2	Pool size: (1, 2, 2), Strides: (1, 2, 2)	(Batch, 75, 11, 35, 256)	0
Conv3D_3	Filters: 75, Kernel: 3×3×3, Padding: 'same'	(Batch, 75, 11, 35, 75)	518,475
Activation	ReLU	(Batch, 75, 11, 35, 75)	0
MaxPool3D_3	Pool size: (1, 2, 2), Strides: (1, 2, 2)	(Batch, 75, 5, 17, 75)	0
TimeDistributed (Flatten)	Flattens spatial dimensions per frame	(Batch, 75, 6375)	0
Bidirectional LSTM_1	Units: 128 (per direction), Return Sequences: True	(Batch, 75, 256)	6,660,096
Dropout_1	Rate: 0.5	(Batch, 75, 256)	0
Bidirectional LSTM_2	Units: 128 (per direction), Return Sequences: True	(Batch, 75, 256)	394,240
Dropout_2	Rate: 0.5	(Batch, 75, 256)	0
TimeDistributed (Dense)	Units: vocab_size + 1 (41), Activation: 'softmax'	(Batch, 75, 41)	10,537


Total Trainable Parameters: 8,471,924 (32.32 MB), as reported by model.summary().

Key Architectural Decisions:

- **Temporal Preservation in Pooling:** The MaxPooling3D layers use a stride of (1, 2, 2), preserving the full 75-frame temporal resolution while progressively downsampling the spatial dimensions. This ensures that the recurrent layers receive a dense temporal sequence.
- **Orthogonal Initialization:** The LSTM recurrent kernels are initialized using the Orthogonal initializer, which helps maintain gradient norm stability in deep recurrent networks.
- **Dropout Regularization:** A dropout rate of 0.5 is applied after each BiLSTM layer to mitigate overfitting, given the limited size of the training dataset (900 samples).

4.3 Algorithm 2: 2D Convolutional Neural Network + Transformer Encoder

The second architecture replaces the recurrent sequence modeling component with a Transformer Encoder. This model explores a more modern paradigm where temporal relationships are modeled entirely through self-attention mechanisms.



4.3.1 Conceptual Background

Two-Dimensional Convolutions for Spatial Encoding: The initial stage of this model employs a shared 2D CNN that processes each of the 75 frames independently. The CNN learns to extract a compact feature vector that summarizes the spatial configuration of the lips and mouth in a single frame. Because the same CNN weights are applied to every frame (via the TimeDistributed wrapper), the model learns translation-invariant visual features.

Transformer Encoder and Self-Attention: The Transformer architecture (Vaswani et al., 2017) eschews recurrence in favor of multi-head self-attention. For each position (query) in the sequence, the self-attention mechanism computes a weighted sum of all positions (keys/values), where the weights are determined by a learned compatibility function. This operation is fully parallelizable and provides each frame with a direct "attention pathway" to every other frame, regardless of their temporal distance. This global receptive field is theoretically advantageous for capturing long-range co-articulation phenomena.

Since self-attention is inherently permutation-invariant, positional encodings must be added to the input embeddings to inject information about the order of frames. In this implementation, the Transformer block uses pre-normalization residual connections for training stability.

4.3.2 Layer-by-Layer Architecture Specification

The Transformer model is constructed using the Keras Functional API to accommodate the custom Transformer block. The following is derived from the notebook implementation.

Shared 2D CNN (Frame-Level Feature Extractor):

Layer (Type)	Configuration	Output Shape (per frame)
Input	-	(46, 140, 1)
Conv2D_1	Filters: 32, Kernel: 3×3, Padding: 'same', ReLU	(46, 140, 32)
MaxPool2D_1	Pool size: (2, 2)	(23, 70, 32)
Conv2D_2	Filters: 64, Kernel: 3×3, Padding: 'same', ReLU	(23, 70, 64)
MaxPool2D_2	Pool size: (2, 2)	(11, 35, 64)
Conv2D_3	Filters: 128, Kernel: 3×3, Padding: 'same', ReLU	(11, 35, 128)
GlobalAveragePooling2D	-	(128)

This CNN is applied to all 75 frames via TimeDistributed, producing an output sequence of shape (Batch, 75, 128). The total parameters for this CNN (including the TimeDistributed wrapper) are **92,672**.

Transformer Encoder Model:

Layer (Type)	Configuration	Output Shape	Parameters
Input	Video tensor	(Batch, 75, 46, 140, 1)	0
TimeDistributed(CNN)	As above	(Batch, 128)	75, 92,672
Dense Projection	Units: 128	(Batch, 128)	75, 16,512
TransformerBlock_1	embed_dim=128, num_heads=4, ff_dim=256	(Batch, 128)	75, 132,480

Layer (Type)	Configuration	Output Shape	Parameters
TransformerBlock_2	embed_dim=128, num_heads=4, ff_dim=256	(Batch, 75, 128)	132,480
TimeDistributed(Dense)	Units: 41, Softmax	(Batch, 75, 41)	5,289

Total Trainable Parameters: 379,433 (1.45 MB), as reported by model.summary().

4.3.3 Transformer Block Implementation Detail

The custom TransformerBlock layer, as defined in the notebook, implements the core self-attention operation with a pre-normalization (pre-norm) residual structure:

```
class TransformerBlock(tf.keras.layers.Layer):
```

```
    def __init__(self, embed_dim, num_heads, ff_dim, dropout=0.1, **kwargs):
```

```
        super().__init__(**kwargs)
```

```
        self.attn = tf.keras.layers.MultiHeadAttention(
```

```
            num_heads=num_heads, key_dim=embed_dim // num_heads
```

```
        )
```

```
        self.ffn = tf.keras.Sequential([
```

```
            tf.keras.layers.Dense(ff_dim, activation='relu'),
```

```
            tf.keras.layers.Dense(embed_dim)
```

```
        ])
```

```
        self.norm1 = tf.keras.layers.LayerNormalization(epsilon=1e-6)
```

```
        self.norm2 = tf.keras.layers.LayerNormalization(epsilon=1e-6)
```

```
        self.drop1 = tf.keras.layers.Dropout(dropout)
```

```
        self.drop2 = tf.keras.layers.Dropout(dropout)
```

```
    def call(self, x, training=False):
```

```
        # Pre-norm self-attention with residual connection
```

```
        x = self.norm1(x + self.drop1(self.attn(x, x), training=training))
```

```
        # Pre-norm feed-forward with residual connection
```


```
        x = self.norm2(x + self.drop2(self.ffn(x), training=training))
```

return x

The pre-norm configuration (applying layer normalization before the attention and feed-forward sub-layers) has been shown to improve training stability and convergence in Transformer models.

4.4 Connectionist Temporal Classification (CTC)

Both models are trained using **Connectionist Temporal Classification (CTC)** loss (Graves et al., 2006). CTC is a specialized objective function designed for sequence-to-sequence problems where the alignment between input timesteps and output labels is unknown a priori.



How CTC Works:

The model outputs a probability distribution over the vocabulary (including a special "blank" token) for each of the 75 timesteps. CTC considers all possible valid alignments of length 75 that collapse to the target transcription. The collapsing rules are:

1. Merge consecutive repeated characters.
2. Remove all blank tokens.

For example, the alignments $b\emptyset\emptyset ii\emptyset n\emptyset$ and $\emptyset b\emptyset i\emptyset\emptyset n\emptyset$ both map to the string "bin". CTC computes the total probability of the target transcription by summing the probabilities of all such valid alignments. The loss is the negative log-likelihood of this total probability. During training, the network learns to increase the probability mass on frames that correspond to character emissions while assigning blank tokens to silent or ambiguous frames.

CTC Decoding During Inference:

During evaluation, a **greedy CTC decoder** is employed to convert the per-frame probability matrix into a text string:

1. For each timestep, select the character index with the highest probability (argmax).
2. Merge consecutive duplicate indices.
3. Remove all blank token indices.
4. Map the remaining integer sequence back to characters using the inverse vocabulary lookup.

The implementation in the notebook uses `tf.keras.backend.ctc_decode` with `greedy=True` for this purpose.

5. Implementation Details

This section documents the technical environment, tools, libraries, and specific configuration parameters used to realize the project.

5.1 Development Environment and Libraries

Component	Specification
Platform	Kaggle Cloud Notebooks
GPU Acceleration	2 × NVIDIA T4 (16 GB VRAM each)
Deep Learning Framework	TensorFlow 2.19.0 with Keras API
Programming Language	Python 3.12
Video I/O	opencv-python-headless
Array Operations	NumPy
Visualization	Matplotlib, imageio (for GIF creation)
Data Download	gdown (Google Drive downloader)

GPU memory growth is configured at startup using `tf.config.experimental.set_memory_growth` to allow incremental allocation, preventing the framework from pre-allocating the entire 16 GB of VRAM and ensuring stable operation with dual GPUs.

5.2 Vocabulary Construction and Encoding

The character-level vocabulary is defined as follows:

Characters: a b c d e f g h i j k l m n o p q r s t u v w x y z 1 2 3 4 5 6 7 8 9 ' ! ? <space>

Vocabulary Size: 40 characters.

Encoding: TensorFlow's `StringLookup` layer maps each character to a unique integer index (0–39). The inverse `StringLookup` layer is used during inference to decode integer sequences back to text.

CTC Blank Token: The blank token is placed at index `VOCAB_SIZE` (index 40), which is outside the range of the regular vocabulary. This prevents any collision between blank emissions and valid character predictions.

5.3 Model Hyperparameters and Training Configuration

The following table summarizes the key hyperparameters and training configuration for both models, as implemented in the notebook.

Hyperparameter / Setting	3D CNN + BiLSTM Model	2D CNN + Transformer Model
Batch Size	2	2
Optimizer	Adam	Adam
Learning Rate	0.0001	0.0001
Loss Function	CTC Loss (blank_index=40)	CTC Loss (blank_index=40)
Input Shape	(75, 46, 140, 1)	(75, 46, 140, 1)
Epochs (Max)	15 (fine-tuning)	20 (from scratch)
Early Stopping	Patience = 5, monitor = val_loss	Patience = 5, monitor = val_loss
Learning Rate Schedule	Exponential decay (factor = $\exp(-0.1)$ after epoch 30)	Exponential decay (factor = $\exp(-0.1)$ after epoch 30)
Model Checkpoint	Save best weights to .h5	Save best weights to .h5
Total Trainable Parameters	8,471,924	379,433

Callbacks:

A suite of Keras callbacks is employed during training:

- **ModelCheckpoint:** Saves the model weights that achieve the lowest validation loss.
- **EarlyStopping:** Halts training if validation loss does not improve for 5 consecutive epochs, preventing overfitting and conserving compute resources.
- **LearningRateScheduler:** Applies an exponential decay to the learning rate, which can help the model settle into a better minimum.
- **ShowPred (Custom Callback):** At the end of each epoch, this callback selects a fixed sample from the validation set, generates a prediction using the current model weights, and prints both the ground truth and the predicted transcript. This provides qualitative, human-interpretable monitoring of training progress.

5.4 Transfer Learning and Weight Initialization

A critical implementation detail for the BiLSTM model was the loading of pre-trained weights from a checkpoint file trained for 96 epochs on the same GRID Corpus Speaker 1 subset.

Challenge: The checkpoint was originally created using an older version of TensorFlow, and the variable naming conventions do not directly match the layer names generated by the Keras API in TensorFlow 2.x.

Solution: A manual weight mapping script was developed, as shown in the notebook:

1. `tf.train.load_checkpoint(checkpoint_path)` is used to read the raw dictionary of variable names and tensor values from the checkpoint file.
2. Through inspection, the variable naming patterns are analyzed (e.g., `layer_with_weights-3/forward_layer/cell/kernel`).
3. Each tensor is manually assigned to the corresponding Keras layer using the `layer.set_weights()` method. This process requires precise matching of tensor shapes and careful alignment of kernel, recurrent kernel, and bias components for LSTM layers.

The loaded epoch-96 weights were saved to `/kaggle/working/models/bilstm/epoch96.weights.h5` for subsequent reuse.

The Transformer model, in contrast, is trained **from scratch** with random weight initialization (default Glorot uniform for dense layers).

6. Results & Evaluation

Both trained models were evaluated on the 100-sample held-out test set using standard sequence error metrics. The training dynamics were analyzed through loss curves, and qualitative predictions were examined.

6.1 Quantitative Performance Metrics

The primary evaluation metrics are **Character Error Rate (CER)** and **Word Error Rate (WER)**. Both are derived from the Levenshtein edit distance, which counts the minimum number of insertions, deletions, and substitutions required to transform the predicted string into the reference string. The error rate is the edit distance divided by the length of the reference sequence.

The evaluation was performed using a custom evaluate function that iterates over the test dataset, decodes predictions using greedy CTC, and computes CER and WER for each sample. The final reported values are averages over the entire test set of 100 samples.

Model	Training Regimen		Character Error Rate (CER) ↓	Word Error Rate (WER) ↓
3D CNN + BiLSTM	Fine-tuned for 15 epochs (from epoch 96 checkpoint)		0.6224	0.8833
2D CNN + Transformer	Trained from scratch for 20 epochs		0.8125	1.0000




Figure 3: Model Comparison on Test Set — CER and WER Bar Chart

Figure Description: Bar chart comparing the Character Error Rate (CER) and Word Error Rate (WER) of both models on the 100-sample held-out test set. The blue bars represent the 3D CNN + BiLSTM model (CER: 0.622, WER: 0.883). The green bars represent the CNN + Transformer model (CER: 0.812, WER: 1.000). Lower values indicate better performance. The BiLSTM model demonstrates superior accuracy on both metrics, attributable to transfer learning from pre-trained weights and the data efficiency of recurrent architectures.

Interpretation:

The 3D CNN + BiLSTM model achieves a CER of 62.24%, meaning that, on average, about 62 characters per 100 are incorrect. While this error rate is substantial, it represents a non-trivial learning signal given the inherent difficulty of the task and the limited dataset size. The corresponding WER of 88.33% indicates that most predicted sentences contain at least one word error.

The Transformer model exhibits higher error rates (CER: 81.25%, WER: 100.00%). This outcome is expected and consistent with the known characteristics of attention-based models. Transformers typically require larger datasets and longer training durations to surpass recurrent architectures. With only 900 training samples and 20 epochs, the Transformer has not fully converged. However, the fact that the model achieves a CER below 1.0 demonstrates that it has learned some underlying structure from the data.

6.2 Training and Validation Loss Curves

The training and validation loss were logged for each epoch. The following table presents the recorded values from the notebook execution.

BiLSTM Loss Values (15 Epochs):

Epoch	Train Loss	Val Loss
1	70.04	32.96

Epoch	Train Loss	Val Loss
2	42.04	30.57
3	34.51	24.01
4	27.05	18.94
5	22.69	13.03
6	17.84	11.43
7	14.89	10.36
8	12.86	8.16
9	11.68	6.21
10	9.93	6.69
11	8.79	5.98
12	8.36	5.63
13	7.88	4.89
14	9.16	5.08
15	7.30	4.21

Transformer Loss Values (First 15 of 20 Epochs Shown):

Epoch	Train Loss	Val Loss
1	123.95	108.08
2	106.74	97.31
3	96.88	89.70
4	90.86	88.61
5	87.75	85.72
6	85.59	84.76
7	83.23	81.90

Epoch	Train Loss	Val Loss
8	82.34	81.47
9	81.21	80.76
10	80.03	78.54
11	79.08	77.98
12	78.22	77.12
13	77.13	76.08
14	76.79	75.35
15	76.30	73.80
...	...	...
20	72.84	70.26

Analysis of Loss Dynamics:

- BiLSTM Model:** The model exhibits rapid convergence. The significant drop in loss from epoch 1 to epoch 15 is attributable to the strong initialization provided by the pre-trained checkpoint. The model essentially continues an already well-advanced optimization trajectory. The narrow gap between training and validation loss throughout training indicates that the model is generalizing well and not overfitting to the training set, a benefit of the Dropout regularization applied.
- Transformer Model:** The model starts with a much higher loss due to random initialization. However, it demonstrates a consistent and monotonic downward trend in both training and validation loss. Crucially, the validation loss tracks the training loss closely and shows no signs of divergence, which would indicate overfitting. The learning curves suggest that the model is on a stable path to convergence and would likely achieve significantly lower error rates if training were extended for additional epochs (e.g., 100–200 epochs).




Figure: Training and Validation Loss Curves for Both Models

Figure Description: Side-by-side comparison of CTC loss trajectories over training epochs. The left panel shows the 3D CNN + BiLSTM model, which exhibits rapid convergence from the pre-trained initialization (final train loss: 7.30, val loss: 4.21). The right panel shows the CNN + Transformer model, which demonstrates stable, monotonic learning from scratch (final train loss: 72.84, val loss: 70.26).

6.3 Qualitative Analysis of Sample Predictions

The notebook includes a "Live Demo" cell (Cell 16) that runs both models on several test videos and prints the predictions alongside the ground truth. The following table presents a selection of these predictions to provide a more intuitive understanding of model behavior.

Video File	Ground Truth	BiLSTM Prediction	BiLSTM CER	Transformer Prediction	Transformer CER
bbaf2n.mpg	"bin blue at f two now"	"sen blue bn si soon"	0.524	"b "	0.857
bbaf3s.mpg	"bin blue at f three soon"	"sen re in one sown"	0.583	"b e l e l "	0.708
bbaf4p.mpg	"bin blue at f four please"	"bin blue n n ninen sooon"	0.560	"bl a "	0.760
bbaf5a.mpg	"bin blue at f five again"	"bin gree in s sive sgain"	0.333	"sl "	0.917
bbal6n.mpg	"bin blue at l six now"	"bin bree in sinee again"	0.667	"la "	0.810
bbal7s.mpg	"bin blue at l seven soon"	"bin blue in s nie sooon"	0.333	"be e e"	0.792

Observations:

- **BiLSTM Predictions:** The BiLSTM model often captures the approximate structure of the sentence, correctly identifying several words or partial words. Common errors include confusion between visually similar letters (e.g., 'f' and 's', 'l' and 'r') and repetition of characters (e.g., "sooon"). These errors are consistent with the fundamental challenge of viseme ambiguity.
- **Transformer Predictions:** The Transformer model's predictions are significantly shorter and often consist of only a few characters or repeated letters. This behavior is typical of a Transformer in early training stages on limited data; the model has learned to output high-probability tokens but struggles to generate coherent, full-length sequences. The high CER values reflect this tendency to collapse to very short outputs.

These qualitative examples reinforce the quantitative metrics and highlight the need for more training data and epochs for the Transformer model to achieve competitive performance.

7. Comparative Analysis

This section provides a detailed, multi-faceted comparison of the two implemented algorithms, synthesizing the quantitative results, architectural properties, and practical considerations.

7.1 Architectural Trade-offs

Aspect	3D CNN + BiLSTM	2D CNN + Transformer
Spatial Feature Extraction	3D Convolutions capture joint spatial and temporal motion patterns directly.	2D Convolutions extract per-frame spatial features; motion must be inferred by the subsequent temporal model.
Temporal Modeling	Bidirectional LSTM processes sequence step-by-step, maintaining a hidden state that summarizes past context.	Transformer Encoder with multi-head self-attention models global dependencies; each frame can attend directly to any other frame.
Long-Range Dependencies	Handled via gated memory cells (LSTM). Effective for sequences of moderate length (75 frames).	Handled via direct attention pathways ($O(1)$ path length between any two frames). Theoretically superior for very long sequences.
Parallelization	Inherently sequential. The computation for timestep t depends on the hidden state from timestep $t-1$.	Fully parallelizable across the time dimension. All frames can be processed simultaneously in the self-attention layer.
Data Efficiency	High. Recurrent inductive biases allow the model to learn effectively from relatively small datasets.	Low. Requires larger datasets or extensive pre-training to learn sequential structure from scratch.
Interpretability	Hidden states and gates can be analyzed to understand information flow over time.	Attention weights provide a direct visualization of which frames the model is "looking at" when making a prediction.
Model Size	~8.5 million parameters.	~0.38 million parameters.

7.2 Convergence Behavior and Data Efficiency

The results of this project starkly illustrate the difference in data efficiency between the two architectures.

- **BiLSTM with Transfer Learning:** The BiLSTM model, initialized with features learned over 96 prior epochs, converged to a validation loss of 4.21 after just 15 additional epochs. This demonstrates the power of transfer learning and the sample efficiency of recurrent networks for this task. The model effectively leveraged prior knowledge to achieve a usable level of performance with minimal additional computation.
- **Transformer from Scratch:** The Transformer model, despite having an order of magnitude fewer parameters, required 20 epochs to reach a validation loss of 70.26. While the loss curve is healthy and descending, it underscores that Transformers are not "plug-and-play" replacements for RNNs in low-data regimes. Their strength lies in scalability—with sufficient data (e.g., the full 34-speaker GRID Corpus), they are expected to outperform RNNs.

7.3 Computational Characteristics

Characteristic	3D CNN + BiLSTM	2D CNN + Transformer
Training Step Time	Slower per batch due to sequential LSTM computation.	Faster per batch due to parallelized attention.
GPU Memory Usage	High, primarily due to 3D convolution activations.	Moderate; the attention matrix (75×75) is relatively small.
Scalability	Limited by sequential bottleneck.	Excellent; scales efficiently with additional compute and data.

Criterion	BiLSTM	Transformer
Parameters	8,471,924	379,433
Epochs trained	96+30	50
Spatial modeling	3D CNN	2D CNN
Temporal modeling	BiLSTM	Self-Attention
Parallelizable	No	Yes
CTC Loss	Yes	Yes
CER (test)	0.6224	0.8125
WER (test)	0.8833	1.0000

Summary of Comparative Analysis:

For the specific constraints of this project (900 training samples, 15–20 epochs), the 3D CNN + BiLSTM model is the superior performer. It benefits significantly from transfer learning and the inductive bias of recurrence. However, the 2D CNN + Transformer model represents the more scalable and forward-looking architecture. Its performance on this small dataset is an expected consequence of training from scratch, not a fundamental flaw. Given the trajectory of the field, future state-of-the-art VSR systems will almost certainly be built upon Transformer or hybrid attention-based foundations.

8. Conclusion

This project successfully designed, implemented, and critically evaluated two distinct deep learning architectures for end-to-end visual speech recognition on the GRID Corpus. The entire pipeline—from raw video preprocessing and custom data pipeline engineering to model construction, transfer learning, and rigorous evaluation—was executed within the Kaggle cloud environment using TensorFlow 2.19.0 and dual NVIDIA T4 GPUs.

Key Findings:

1. **Effective Transfer Learning:** The 3D CNN + BiLSTM model, fine-tuned from a pre-trained checkpoint, achieved a **Character Error Rate (CER) of 62.24%** and a **Word Error Rate (WER) of 88.33%**. This result validates the architectural combination of spatiotemporal convolutions and bidirectional recurrence for lip reading. The rapid convergence and stable training dynamics demonstrate that transfer learning is a highly effective strategy for this domain, especially when data is limited.
2. **Transformer Viability and Scalability:** The 2D CNN + Transformer model, trained from scratch, achieved a CER of 81.25%. While the absolute performance lags behind the fine-tuned BiLSTM, the model exhibited stable, convergent training behavior. This confirms that the Transformer architecture is fundamentally sound for VSR and would benefit substantially from more data and extended training. Its parallelizable nature and global receptive field position it as the more promising long-term architecture.
3. **Architectural Trade-offs Confirmed:** The project provided empirical evidence for well-understood theoretical trade-offs. Recurrent models with transfer learning excel in low-data, rapid-deployment scenarios. Transformer models, while data-hungry, offer superior scalability and are the preferred choice when large-scale training resources are available.
4. **Reproducible Research Framework:** The project delivers a fully documented, self-contained Jupyter Notebook. All code—including the custom CTC loss implementation, manual checkpoint weight loading, Transformer block definition, and evaluation metrics—is clearly presented and executable, providing a robust foundation for future work in this area.

9. Future Work

The current project establishes a strong baseline and opens several promising avenues for future research and development. The following enhancements are identified as high-priority next steps:

1. **Full Corpus Scaling:** The most impactful immediate improvement would be to extend training to the **entire GRID Corpus** (all 34 speakers, approximately 34,000 video samples). This volume

of data would allow the Transformer model to fully converge and is expected to yield a CER substantially lower than the BiLSTM baseline, demonstrating the scalability advantage of attention mechanisms.

2. **Advanced Decoding with Language Models:** The current greedy CTC decoder operates purely at the character level. Implementing a **beam search decoder** integrated with an n-gram language model (trained on the GRID grammar) would incorporate syntactic constraints. This would likely reduce the WER dramatically by penalizing grammatically invalid word sequences and resolving homophone ambiguities.
3. **Robust Face and Landmark Detection:** The static mouth crop [190:236, 80:220] is specific to the GRID Corpus recording setup. To generalize to real-world, in-the-wild videos, this preprocessing step should be replaced with a dynamic facial landmark detector, such as **MediaPipe Face Mesh** or **dlib**. This would enable the system to handle varying head poses, distances, and speaker morphologies.
4. **Multimodal Audio-Visual Fusion:** Robust speech recognition in challenging environments often relies on fusing audio and visual streams. A logical extension of this project is to incorporate an acoustic branch (e.g., processing Mel-Frequency Cepstral Coefficients or spectrograms) that fuses with the visual stream. Such a multimodal model would be resilient to both acoustic noise and visual occlusions, representing the current frontier in ASR/VSR research.
5. **Leveraging Large-Scale Pre-trained Models:** Instead of training Transformers from scratch, future work should explore fine-tuning large, self-supervised audio-visual models such as **AV-HuBERT**. These models are pre-trained on thousands of hours of unlabeled video and have already learned rich, generalizable visual speech representations. Fine-tuning such a model on the GRID Corpus would likely achieve state-of-the-art results with minimal additional training.
6. **Real-Time Inference Optimization:** To move toward practical deployment, the model could be optimized for real-time inference. Techniques such as quantization (reducing weight precision to INT8) and pruning could reduce the model's memory footprint and latency, enabling a browser-based or mobile lip-reading assistant.

10. References

1. Assael, Y., Shillingford, B., Whiteson, S., & de Freitas, N. (2016). *End-to-End Sentence-level Lipreading*. arXiv preprint arXiv:1611.01599. <https://arxiv.org/abs/1611.01599>
2. Cooke, M., Barker, J., Cunningham, S., & Shao, X. (2006). An Audio-Visual Corpus for Speech Perception and Automatic Speech Recognition. *Journal of the Acoustical Society of America*, 120(5), 2421–2424.
3. Vaswani, A., Shazeer, N., Parmar, N., Uszkoreit, J., Jones, L., Gomez, A. N., Kaiser, L., & Polosukhin, I. (2017). Attention Is All You Need. *Advances in Neural Information Processing Systems (NeurIPS)*, 30.
4. Hochreiter, S., & Schmidhuber, J. (1997). Long Short-Term Memory. *Neural Computation*, 9(8), 1735–1780.

5. Graves, A., Fernandez, S., Gomez, F., & Schmidhuber, J. (2006). Connectionist Temporal Classification: Labelling Unsegmented Sequence Data with Recurrent Neural Networks. *Proceedings of the 23rd International Conference on Machine Learning (ICML)*.
6. Ma, P., Haliassos, A., Fernandez-Lopez, A., Chen, H., Petridis, S., & Pantic, M. (2023). Auto-AVSR: Audio-Visual Speech Recognition with Automatic Labels. *Proceedings of the IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP)*.
7. TensorFlow Team. (2024). *TensorFlow 2.x Documentation*. Google LLC. https://www.tensorflow.org/api_docs
8. Kaggle. (2024). *Kaggle Notebooks — GPU Training Platform*. Alphabet Inc. <https://www.kaggle.com>

Tools and Technologies Employed:

- TensorFlow 2.19.0 (Deep Learning Framework)
- OpenCV (Video Frame Decoding)
- NumPy (Numerical Operations)
- Matplotlib (Visualization)
- imageio (GIF Creation)
- gdown (Dataset Download)
- Python 3.12
- Kaggle Cloud GPU Environment (Dual NVIDIA T4)